

Roll no:215229109

Exercise 02a: Map Reduce applications for Word Counting

Previous exercise described how to save input file in to HDFS. This exercise train students to do MapReduce process using word counting application.

Prerequisites

Ensure that Hadoop is installed, configured and is running. More details:

Single Node Setup for first-time users.

Cluster Setup for large, distributed clusters.

MapReduce Overview

Hadoop MapReduce is a software framework for easily writing applications which process vast amounts of data (multi-terabyte data-sets) in-parallel on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.

A MapReduce *job* usually splits the input data-set into independent chunks which are processed by the *map tasks* in a completely parallel manner. The framework sorts the outputs of the maps, which are then input to the *reduce tasks*. Typically both the input and the output of the job are stored in a file-system. The framework takes care of scheduling tasks, monitoring them and re-executes the failed tasks.

Typically the compute nodes and the storage nodes are the same, that is, the MapReduce framework and the Hadoop Distributed File System are running on the same set of nodes. This configuration allows the framework to effectively schedule tasks on the nodes where data is already present, resulting in very high aggregate bandwidth across the cluster.

The MapReduce framework consists of a single master `ResourceManager`, one worker `NodeManager` per cluster-node, and `MRAppMaster` per application.

Minimally, applications specify the input/output locations and supply *map* and *reduce* functions via implementations of appropriate interfaces and/or abstract-classes. These, and other job parameters, comprise the *job configuration*.

The Hadoop *job client* then submits the job (jar/executable etc.) and configuration to the `ResourceManager` which then assumes the responsibility of distributing the

software/configuration to the workers, scheduling tasks and monitoring them, providing status and diagnostic information to the job-client.

Inputs and Outputs

The MapReduce framework operates exclusively on `<key, value>` pairs, that is, the framework

views the input to the job as a set of <key, value> pairs and produces a set of <key, value> pairs as the output of the job, conceivably of different types.

The key and value classes have to be serializable by the framework and hence need to implement the Writable interface. Additionally, the key classes have to implement the WritableComparable interface to facilitate sorting by the framework.

Input and Output types of a MapReduce job:

(input) <k1, v1> -> **map** -> <k2, v2> -> **combine** -> <k2, v2> -> **reduce** -> <k3, v3> (output)

Step 1

Compile WordCount.java and create a jar:


(i) Open Eclipse in Cloudera



(ii) Create 'WordCount' java project




Import following Jar files



(iii)





Step 2

Create following folders in HDFS:

- /input - input directory in HDFS ·
- /output - output directory in HDFS

```
cm_api.py      express-deployment.json  Public  ue
copsysample    hi.txt                 re      Videos
Desktop        kerberos               te      WordCount.jar
Documents      lib                    temp    workspace
Downloads      Music                  TempFile ye
eclipse        parcels                 Templates

[cloudera@quickstart ~]$ hdfs dfs -mkdir /in00
[cloudera@quickstart ~]$ hdfs dfs -ls /
Found 8 items
drwxrwxrwx - hdfs supergroup 0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup 0 2022-08-02 23:05 /hbase
drwxr-xr-x - cloudera supergroup 0 2022-08-03 00:09 /in00
drwxr-xr-x - solr solr 0 2017-10-23 09:18 /solr
drwxr-xr-x - cloudera supergroup 0 2022-07-31 10:24 /temp
drwxrwxrwt - hdfs supergroup 0 2022-07-28 08:53 /tmp
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup 0 2017-10-23 09:17 /var
[cloudera@quickstart ~]$ cat > WCFFile.txt
Big Data is the technique to handle huge amount of data with characteristics of 3v's.
^C
[cloudera@quickstart ~]$ cat

clear
clear

^C
[cloudera@quickstart ~]$ cat WCFFile.txt
Big Data is the technique to handle huge amount of data with characteristics of 3v's.
```

Step 3

Create and copy sample text-files into input folder:

```
[cloudera@quickstart ~]$ hdfs dfs -ls /in00/
```

Found 1 items

```
-rw-r--r-- 1 cloudera supergroup 158 2021-08-15 04:32 /in00/WCFile.txt
```

```
[cloudera@quickstart ~]$ hdfs dfs -copyFromLocal WCFile.txt in00/  
copyFromLocal: `in00/': No such file or directory  
[cloudera@quickstart ~]$ hdfs dfs -copyFromLocal WCFile.txt /in00/  
[cloudera@quickstart ~]$ hdfs dfs -ls /in00/  
Found 1 items  
-rw-r--r-- 1 cloudera supergroup 86 2022-08-03 00:14 /in00/WCFile.txt
```

Step 4

Run the MapReduce application:

```
hadoop jar /home/cloudera/WordCount.jar WordCount /in00/WCFile.txt
```

/out00 Show MapReduce Framework



```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
22/08/03 05:28:55 INFO mapreduce.Job: Running job: job_1658638425302_0002  
22/08/03 05:29:02 INFO mapreduce.Job: Job job_1658638425302_0002 running in uber mode : false  
22/08/03 05:29:02 INFO mapreduce.Job: map 0% reduce 0%  
22/08/03 05:29:11 INFO mapreduce.Job: map 100% reduce 0%  
22/08/03 05:29:17 INFO mapreduce.Job: map 100% reduce 100%  
22/08/03 05:29:18 INFO mapreduce.Job: Job job_1658638425302_0002 completed successfully  
22/08/03 05:29:19 INFO mapreduce.Job: Counters: 49  
File System Counters  
FILE: Number of bytes read=169  
FILE: Number of bytes written=287015  
FILE: Number of read operations=8  
FILE: Number of large read operations=8  
FILE: Number of write operations=8  
HDFS: Number of bytes read=178  
HDFS: Number of bytes written=99  
HDFS: Number of read operations=6  
HDFS: Number of large read operations=8  
HDFS: Number of write operations=2  
Job Counters  
Launched map tasks=1  
Launched reduce tasks=1  
Data-local map tasks=1  
Total time spent by all maps in occupied slots (ms)=6885  
Total time spent by all reduces in occupied slots (ms)=3766  
Total time spent by all map tasks (ms)=6885  
Total time spent by all reduce tasks (ms)=3766  
Total vcore-milliseconds taken by all map tasks=6885  
Total vcore-milliseconds taken by all reduce tasks=3766  
Total megabyte-milliseconds taken by all map tasks=7056240  
Total megabyte-milliseconds taken by all reduce tasks=3856384  
Map-Reduce Framework  
Map input records=4  
Map output records=16  
Map output bytes=131  
Map output materialized bytes=169
```

Step 5

Output:

```
[cloudera@quickstart ~]$ hdfs dfs -ls /out00/
```

Found 2 items

```
-rw-r--r-- 1 cloudera supergroup 0 2021-08-15 04:41 /out00/_SUCCESS -rw-r--r-- 1
```

```
cloudera supergroup 113 2021-08-15 04:41 /out00/part-r-00000
```

```
[cloudera@quickstart ~]$ hdfs dfs -cat /out00/part-r-00000
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Merged Map outputs=1  
GC time elapsed (ms)=75  
CPU time spent (ms)=1340  
Physical memory (bytes) snapshot=492310528  
Virtual memory (bytes) snapshot=3136634888  
Total committed heap usage (bytes)=379068224  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=67  
File Output Format Counters  
Bytes Written=99  
[cloudera@quickstart ~]$ hdfs dfs -cat /out00/part-r-00000  
aaa 1  
bbb 1  
cccc 1  
ddd 1  
eee 1  
fff 1  
ggg 1  
hhhh 1  
iii 1  
jjj 1  
kkk 1  
llll 1  
mmm 1  
nnn 1  
ooo 1  
ppp 1  
[cloudera@quickstart ~]$
```